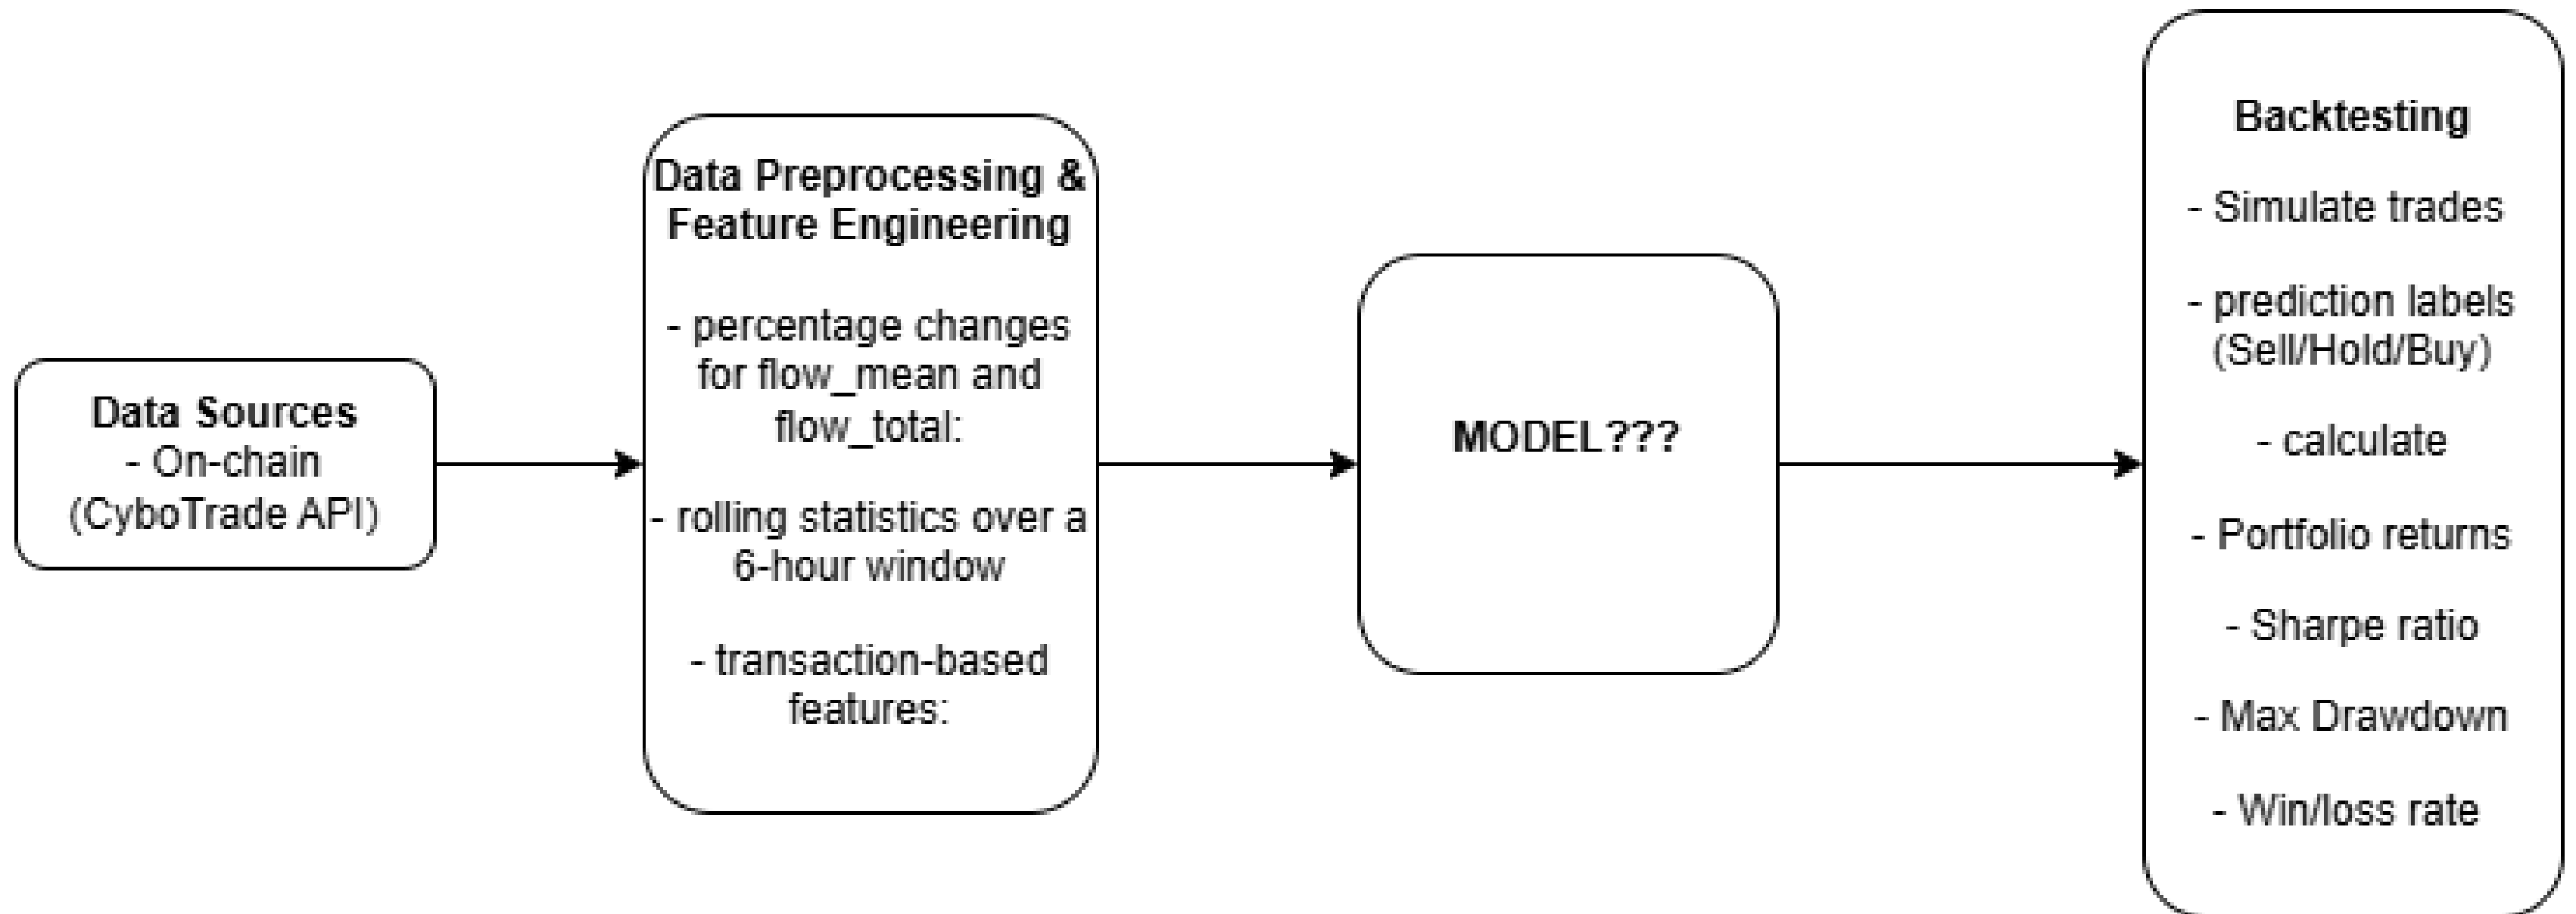




Domain 2

Convolutional Neural Networks for Alpha Strategy Development in Volatile Crypto Markets

Team Arctic Rhinos 1





HMM

Linear Regression / Logistic Regression

K-Means

Transformers

CNN




HMM

HMMs are not the best at modeling complex non-linear patterns, where CNNs, LSTMs, and Transformers perform better.



Linear Regression / Logistic Regression


Assumes fixed relationships between variables, limits ability to adapt to dynamic market conditions that deep learning models handle more flexibly.





K-Means Algo


K-Means assumes fixed clusters and struggles with dynamic patterns, unlike HMMs or LSTMs





Transformers

Require large datasets to perform optimally, making them less effective in data-scarce environments compared to simpler models like Random Forests or Logistic Regression.




CNN

- Learn and extract meaningful patterns from price and indicator sequences without us having to define any state-based assumptions, unlike HMM
- Main goal was to build a system that outputs clear Buy/Hold/Sell predictions that we could plug into a backtesting engine.
- CNNs are inherently built to handle multi-dimensional inputs, while other model might struggle to scale with higher dimensional feature spaces.
- CNNs can be retrained regularly on new data, enabling them to adapt to market regimes and conditions.



CNN Model Building

- Define a CNN architecture
 - Model Compilation
 - Data Preparation
 - Model Training
 - Model Evaluation
 - Confusion Matrix
- 

Some Model Details

- Conv1D layer (64 filters, kernel size 3, ReLU activation) with batch normalization and max pooling.
- Conv1D layer (128 filters, kernel size 3, ReLU activation) with batch normalization and max pooling.
- Flatten layer to convert 2D outputs into 1D.
- Dense layer (64 units, ReLU activation) with dropout for regularization.
- Dense layer with 3 output units (one for each class: Sell, Hold, Buy) and softmax activation.
- Model Compilation:
 - Use Adam optimizer with a learning rate of 0.001.
 - Use categorical crossentropy loss for multi-class classification.
 - Track accuracy as the evaluation metric.

Model Demo